

Lab Book

Brad Day

2025-05-19

Table of contents

Preface	3
1 Reading	4
1.1 Summary of “Broad Iron Lines in Active Galactic Nuclei” [1]	4
2 Initial Gradus Testing	8
3 Parameter Space Investigations	14
3.1 Ruling out α_{22} and α_{52}	14
3.2 Exploring the 5-dimensional parameter space	14
3.2.1 Line profiles	15
3.2.2 Redshift Image Renders	15
3.3 Summary	19
3.4 Code Specifics	19
4 Fitting Spectral Lines	20
5 Fitting to Data	22
5.1 Initial Results	22
6 Developing and Migrating to a Table Model	27
6.1 Introduction	27
6.2 Data Generation	27
6.3 Loading into FITS Format	27
6.4 The Model	28
6.4.1 Performance	28
6.5 Next Steps	29
6.6 Current Issues 2026/07/03	30
7 Notes for the Reduction of XRISM Data	32
7.1 Pipeline Scripts	32
7.1.1 Reprocessing:	32
7.1.2 Analysis Setup	33
7.1.3 PSP Limit Checks	33
7.2 Removing Anomalous Pixels from Xtend Data	33
7.3 Proximity and Rise time screening	37
7.4 Extract Broadband Images and Establish Source Center Coordinates	38
References	40

Preface

This book details my research during Summer 2026. Many thanks to Dr. Andrew Young and Mr Darius Michienzi for supervising this project.

This is a Quarto book.

To learn more about Quarto books visit <https://quarto.org/docs/books>.

1 Reading

1.1 Summary of “Broad Iron Lines in Active Galactic Nuclei” [1]

Fluorescent iron line, $\text{Fe-K}\alpha$, in the X-ray band at $6.4 - 6.9 \text{ keV}$ is the strongest narrow line emitted by an accretion disk around a black hole. Inverse Compton scattering of the soft photons (lower energy than the reactant particles) in a hot corona is the most likely candidate for the production of these X-rays. The emitted line is broadened and skewed due to the Doppler effect and gravitational redshift, and is often called the broad iron line due to being broadened to a FWHM of $5 \times 10^4 \text{ km s}^{-1}$.




Figure 1.1: The profile of the broad iron line, image credit to Fabian (2000) [1].

Soft Comptonisation (multiple inverse Compton scattering by hot thermal electrons) occurs giving a power law X-ray spectrum. Flares from active or flaring regions irradiate the accretion disk forming a “reflection” component in the spectrum.




Figure 1.2: Reflection spectrum from a Monte Carlo simulation of an illuminated slab. The dashed line is the incident continuum and the solid line shows the reflected spectrum. Image credit Reynolds (1996) [1]

The iron line is produced when one of the two K-shell electrons is ejected following absorption of an X-ray, the threshold for neutral iron is 7.1 keV. The resultant state can decay through the dropping of an L-shell electron into the K-shell, releasing 6.4 keV as either an emission line photon (34%) or an Auger electron (66%). Ionised iron has a photoelectric threshold and the $K\alpha$ of higher energy. The strength of the iron line is measured in terms of its equivalent width.

i Equivalent Width

The equivalent width of a spectral line is the width of the rectangle which has the same area as that between the line and the continuum level.




Figure 1.3: A diagram indicating the equivalent width. The width of the rectangle is the equivalent width.

The equivalent width is a function of the geometry of the accretion disk, the elemental abundances, the inclination angle, and the ionisation state.

! Key Points

- **X-ray Continuum:** The hard X-ray continuum in AGNs and GBHCs is thought to be from active/flaring regions in the corona. Thermal electrons multiply inverse Compton scatter optical and UV photons from the disk to X-ray energies. A hard X-ray power law results and irradiates the accretion disk to produce a reflection component.
- **Reflection Component:** The reflection component causes the spectrum to flatten above 10 keV as Compton recoil reduces the backscattered flux and results in a strong fluorescent iron line at ~ 6.4 keV. The energy and strength of this line depends on iron abundance, inclination of the disk, and its ionisation state. A strong iron absorption edge may also be present from an ionised disk.
- **Iron Line Profile:** Doppler shifts, relativistic broadening and gravitational redshifting give the broad and skewed line profile. As the line comes from the innermost regions, these effects are very pronounced and may be used to determine the black hole spin and inclination angle of the accretion disk. In most Seyfert 1 galaxies, the disk is inclined at around 30° , consistent with the standard model.
- **Observations:** The iron line was first clearly detected by Ginga and a line profile subsequently resolved by ASCA confirmed the broad and skewed shape expected from an accretion disk around a Schwarzschild black hole. RXTE has been used to study the line and continuum variability on shorter timescales, with reduced energy resolution.
- **Black Hole Spin:** The radius of the ISCO decreases with spin. Since the line profile is sensitive to the innermost radius of fluorescent emission this may be used to estimate the black hole, with some astrophysical assumptions. With present time-averaged observations such measurements may be ambiguous as alternative models with different values for the spin may produce identical line profiles.
- **Alternative Models for the Production of the Broad Iron Line:** Alternative models that do not require an accretion disk appear to fail, in particular the line width cannot be entirely due to Comptonisation. Hybrid models in which the Comptonisation and Doppler/gravitational effects produce the profile are heavily constrained.
- **Variability:** Rapid X-ray continuum variability is observed in most AGNs, and the iron line is expected to vary in response to this with short time lag. Whilst these timescales are too short to be probed presently, significant and complex iron line variability has been observed. Often the line flux is seen to remain constant whilst the continuum changes, and there appears to be an anticorrelation between the reflected fraction and the equivalent width. In another study the reflected fraction and the photon index of the power law are correlated for individual objects and between different ones. These need to be explained as they appear contrary to our simple model of reflection. Flux correlated changes in the ionisation state of the disk may explain some of these facts.
- **Reverberation Mapping:** The rapid variability is associated with activation of new flares in the corona above the accretion disk. X-ray reverberation mapping uses observations of the iron line response to sudden changes in the continuum to study the accretion disk and black hole. This may be used to determine the geometry of the X-ray emission and the black hole spin and mass.
- **Other Classes of Object** Iron lines are observed in quasars, LLAGNs, radio-loud AGNs, and GBHCs. Quasars have the iron line decrease with increasing luminosity which may be due to the more luminous sources accreting closer to the Eddington limit being more highly ionised. Observations in LLAGNs may determine whether the accretion with low rates is an advection-dominated accretion flow or a thin disk. Weak lines have been observed suggesting that their low accretion rate flows are thin disks rather than optically thick ADAFs. The broad iron lines and reflection humps are weak or absent perhaps because the reflection signature is swamped by a beamed continuum. GBHCs show smeared absorption and emission features in the reflection spectra, and there is a debate as to the precise nature of the accretion flow within a few tens of Schwarzschild radii.

2 Initial Gradus Testing

The below details initial experimentation with `Gradus.jl`. The code block below renders an image of a black hole with a thick accretion disk. Colour denotes the redshift of a given region on the disk.

```
using Gradus
using Plots

# Function defining the cross section of the disk
function discCrossSection(x)
    centre = 8
    radius = 3

    if (x < centre - radius) || (radius + centre < x)
        zero(x)
    else
        r = x - centre
        0.5(sqrt(radius^2 - r^2) + (0.5sin(3x)))
    end
end

# Metric Parameters
M = 1.0      # Mass
a = 0.7      # Black hole spin
13 = 0.0
22 = 0.0
52 = 0.0
3 = 0.0
i = deg2rad(70) # Inclination

# Instantiating the metric
m = JohannsenMetric(M, a, 13, 22, 52, 3)

# Position of the observer
# (x[1], Distance, Inclination, x[4])
x = SVector(0.0, 1000.0, i, 0.0)

# Maximum affine time, parameterises the solution
# _max ~ 2*x[2]
_max = 2000.0

# Setting up impact parameters
= range(-10.0, 10.0, 100)
= range(-10.0, 10.0, 100)
```



```

# Getting an array of velocities from the impact parameters
vs = vec([map_impact_parameters(m, x, a, b) for a in , b in ])
xs = fill(x, size(vs))

# Generating solutions
sols = tracegeodesics(m, xs, vs, _max, save_on = false)

# Generating points and reshaping to be the same shape as the image
points = unpack_solution.(sols.u)
points = reshape(points, (100, 100))

# Redshift point function
redshift = ConstPointFunctions.redshift(m, x)
redshiftGeometry = redshift ConstPointFunctions.filter_intersected()

# Disc
d = ThickDisc(r -> discCrossSection(r))

# this function returns the impact parameter axes
, , image = rendergeodesics(
    m, x, d, _max, pf = redshiftGeometry,
    # image parameters
    image_width = 200, image_height = 150,
    # the "zoom" -- use the impact parameter axes
    lims = (-20, 20), lims = (-15, 15), verbose = false
)

heatmap( , , image, aspect_ratio = 1)

```




Figure 2.1: The first render produced.

Once familiarity was gained with the basics of `Gradus.jl` the line profile functionality was utilised through the following function.

```
function computeLineProfile(m, i, height, bins = range(0.0, 1.5, 180))
    # Position of the observer
    # (x[1], Distance, Inclination, x[4])
    x = SVector(0.0, 10000.0, i, 0.0)

    # Disk
    d = ThinDisc(0.0, Inf)

    # Setting up the model and profile
    model = LampPostModel(h = height)
    profile = emissivity_profile(m, d, model)

    # Computing the line profile
    bins, flux = lineprofile(m, x, d, profile; bins=bins, verbose=true)

    return bins, flux
end
```

This will allow for the investigation of the free parameters for the Johannsen metric; α_{13} , α_{22} , α_{52} , and ϵ_3 . A lamppost corona model is used with variable height h . An example of the effect of varying the height of this corona is shown below.





Figure 2.2: Initial variations of the height of the lamppost corona.

Following initial experimentation the two functions above were integrated together for efficient investigation. The two free parameters that will be the focus of this investigation are α_{13} and ϵ_3 . Examples of the effects of these parameters on the emission lines are shown below






Figure 2.3: Initial variations of α_{13} and ϵ_3

It can be seen that α_{13} has a very smooth, almost linear effect on the shape of the line, while ϵ_3 is much more chaotic, particularly for $\epsilon_3 > 25$.

3 Parameter Space Investigations

Code was set up to easily run several variations to the model which will allow for the investigation of the parameter space. A thin disk accretion model and a lamppost corona were used. It was set up such that for a given configuration of h , a , i , ϵ_3 , and α_{13} the line profile would be calculated and an image of the accretion disk produced and plotted together.

3.1 Ruling out α_{22} and α_{52}

Before exploring the 5D parameter space, tests were ran on both α_{22} and α_{52} such that their effects could be noted, should they be useful in the future.




Figure 3.1: The resultant line profile from varying α_{52} between 0 and 40. A value of 50 did not successfully compute suggesting an upper limit on this parameter.

Varying α_{52} showed near random, and near zero effect on the line and will thus be disregarded for this investigation. Varying α_{22} often failed to compute and will thus similarly be disregarded for this investigation.

3.2 Exploring the 5-dimensional parameter space

Before varying the parameters it was necessary to determine the limits for useful variations. The spin is fundamentally limited to be $|a| < |M|$ and will thus be varied between 0 and 1. The ranges and default values for the parameter space are tabulated in Table 3.1.

Table 3.1: The ranges and default values for the parameter space.

Parameter	Minimum	Maximum	Default
α_{13}	0.0	50.0	0.0
ϵ_3	0.0	30.0	0.0
a	0.0	1.0	0.998
h	R_{ISCO}	$30 R_G$	$10 R_G$
i	0.0°	90.0°	60°

3.2.1 Line profiles

- Figure 3.2 shows that increasing the free parameter α_{13} sharpens the two peaks of the line, increases the maximum flux, and makes it more narrow.
- Figure 3.3 shows that increasing the free parameter ϵ_3 similarly increases the maximum flux, however the resultant profile is much more noisy. While α_{13} provided sharp peaks, ϵ_3 provided broader peaks, particularly in the shorter one.
- Figure 3.4 shows that increasing the inclination smoothes out the line profile and shifts the peak to higher energies. The line is broadened significantly.
- Figure 3.5 shows that increasing the corona height makes the line profile sharper, increases the maximum flux, and shifts the maximum to slightly lower energies.
- Figure 3.6 shows that the spin had the least prominent impact on the profile of the line, seeming to only slightly decrease the height of the peaks with increasing spin. This may be an effect of the small range of spins available. It may therefore be informative to investigate this further by increasing the mass, using different corona heights, or varying the inclination.

3.2.2 Redshift Image Renders

Alongside the line profiles, images of the accretion disk were rendered with a colourmap denoting the redshift of light from a given part of the disk. The corona height was not varied here as it did not impact the render.

Some of the parameters appeared to change the radius of the ISCO. This was investigated by explicitly finding the ISCO, as tabulated below.

Parameter	Value	ISCO (R_G)
α_{13}	0.00	1.24
	16.67	8.62
	33.33	11.70
	50.00	14.03
ϵ_3	0.00	1.24
	10.00	2.36
	20.00	2.90
	30.00	3.51
a	0.00	6.00
	0.32	4.92
	0.63	3.68
	0.95	1.94

The ISCO is proportional to both α_{13} and ϵ_3 , and inversely proportional to the spin.




Figure 3.2: The line profiles resultant from varying α_{13} between 0 and 50.




Figure 3.3: The line profiles resultant from varying ϵ_3 between 0 and 30. Values of 40 and 50 were tested however caused an error, there must be an upper bound on the allowed values of this parameter.




Figure 3.4: The line profiles resultant from varying the inclination between 5 and 85 degrees.




Figure 3.5: The line profiles resultant from varying h between 5 and 30 gravitational radii.




Figure 3.6: The line profiles resultant from varying a between 5 and 95% of $M(=1)$.




Figure 3.7: The effect of α_{13} on the accretion disk render. The ISCO appears to grow with α_{13}




Figure 3.8: The effect of ϵ_3 on the redshift image. The shape of the disk does not change but the redshift profile does, particularly in the shrinking of the blueshift region, and an overall decrease in the degree of redshifting.




Figure 3.9: The effect of inclination on the redshift image. The redshift becomes more apparent as the inclination increases and the disk becomes edge on. This agrees with the observed line profile trend of energy increasing with inclination as the effects of doppler shifting become more apparent with increasing energy.




Figure 3.10: The effect of black hole spin on the redshift image. The ISCO appears to shrink.

3.3 Summary

! Summary

- α_{22} and α_{52} had unpredictable or minimal effect on the line profile and rendered images so are therefore to be disregarded.
- a : Increasing the spin appears to shrink the ISCO of the accretion disk. Its effect on the line profile is less obvious, it appears to decrease the flux slightly but is otherwise negligible.
- h : Increasing the corona height makes the line profile sharper and increases the peak flux, while decreasing the energy of the peak. It has no effect on the renders.
- θ : Increasing the inclination (where 0 degrees has the accretion disk face-on) smoothes out the line profile and increases the peak energy. The effect of redshift becomes more apparent with increased inclination.
- α_{13} : Increasing α_{13} appears to increase the radius of the ISCO and make the line profile sharper.
- ϵ_3 : Increasing ϵ_3 increases the height of the line profile, and makes it more chaotic.
- The ISCO is proportional to both α_{13} and ϵ_3 , and inversely proportional to the spin.

3.4 Code Specifics

The workhorse of the parameter investigations was the utility function `JohannsenParamVar` which sets up the metric and observer position and then calls either `ComputeLineProfile` or `RenderImage`. This utility function was made versatile using keyword arguments allowing for different methods to be called within it. This was then wrapped in `ParamLoop` which took the index (the variable to be looped through) and a set of values to call `JohannsenParamVar` for.

To maximise readability of the redshift images, a custom colormap is made through `RedshiftColormap` which gives red for any values < 1 , blue for those > 1 , and white for those $= 1$. This gave the images seen in Figure [3.7](#), [3.8](#), [3.9](#), [3.10](#).

4 Fitting Spectral Lines

With the parameter space explored, it was useful to setup the architecture for fitting the model to spectral data. This was done using `SpectralFitting.jl` and defining a custom model using `Gradus` which fits the following parameters:

Parameter	Meaning
K	Normalisation
h	Corona Height
E	Line Energy
R_in	Inner Radius
R_out	Outer Radius
	Inclination
a	Spin
13	α_{13} Parameter
3	ϵ_3 Parameter

The architecture of this model was heavily based upon previous work on [GitHub](#).

The model is instantiated using the `LampPostJohannsen` utility function. When called, the parameters are used to instantiate a `JohannsenMetric`, observer position, `ThinDisc`, and an emissivity profile. These are then used to compute the line profile as previously done which is returned.

All fitting is handled by `SpectralFitting` through the Levenberg Marquadt algorithm.

To test the implementation a line profile was computed and random noise added to it. This was then passed through the fitting function. The fitting process currently takes a long time and yields poor results so more work is needed. Time was taken to streamline the functionality of `main2.jl` via keyword arguments and general code cleanup.

Upon fitting, the error `Warning: No parameter uncertainty estimation due to error: LinearAlgebra.Singular` occurs and therefore the fit is not currently quantified.

With the current errors present it was decided to attempt the fit on actual data in order to determine whether they are a result of the data, or the fitting algorithm.




Figure 4.1: The first attempt at fitting to the noisy simulated data. This fit took 67 iterations and 15.3 minutes to complete.

5 Fitting to Data

Data from the NuSTAR (Nuclear Spectroscopic Telescope Array) spacecraft [2] was used, and the dataset used for testing was nu80402315002B01. Data loading was handled by `SpectralFitting.jl` through the `OGIPDataset` struct. This takes the spectrum file as an argument and automatically pulls in the background, `.arf`, and `.rmf` files to compute the spectrum.

5.1 Initial Results

Table 5.1: The ranges and default values for the parameter space.

Model	Variable	Value	Error	χ^2
Johannsen	a	0.57	1.88	116.84
	h	50.00	8.77	
	θ	58.24	4.93	
	α_{13}	2.29	10.96	
	ϵ_3	29.74	25.89	
	K	161.8	7.15	
	E	6.27	0.016	
Kerr	a	0.77	0.057	124.64
	h	50.0	2.75	
	θ	55.44	2.67	
	K	162.0	7.07	
	E	6.27	0.014	




Figure 5.1: The first attempt at fitting to data. The Johannsen metric was used to produce the blue line, and the Kerr metric was used to produce the red. Neither fit was particularly good which may either be the result of poor data, or a poor model.

The NuSTAR datasets have two distinct spectra from each telescope, which can be fit simultaneously, binding some of the parameters together. The fitting code was thus adjusted to fit both datasets simultaneously while binding all parameters other than the normalisation. The dataset labelled “nu80402315002” was once again used.

It is evident from Figure 5.2 that the fitting process needs refinement as the fits do not currently fit the data well, and therefore drawing comparisons would be useless.

Following these experiments it was decided to fit a composite model of **PowerLaw + LampPost**, with the latter being the line profiles generated using either the Kerr or Johannsen metric through `Gradus.jl`. The first attempt at this provided the plots below.




Figure 5.2: The fits to the nu80402315002 dataset. The Johannsen fit is shown in blue, and the Kerr fit in red.




Figure 5.3: The composite fit using the Johannsen metric.




Figure 5.4: The composite fit using the Kerr metric

The inner and outer radii of the accretion disk were set constant at the ISCO and $400 R_g$ respectively. The Johannsen metric currently seems to yield a better fit, however this appears to be an effect of parameters other than the deformation parameters α_{13} and ϵ_3 as they were both relatively small (1.0285 and 1.6702×10^{-5} , respectively). It may therefore be more effective to investigate the effect of these parameters by first fitting the Kerr metric to the data to find the Kerr parameters, then fitting the Johannsen metric with the found parameters, only allowing the two deformation parameters to vary.


Further reading suggests that it may be possible to constrain the parameters for fitting further. In T. Johannsen's 2013 paper describing the Johannsen metric [3], the following constraints were derived for the deformation parameters:

$$\epsilon_3 > -\frac{(M + \sqrt{M^2 - a^2})^3}{M^3} \quad (5.1)$$

and

$$\alpha_{13} > -\frac{(M + \sqrt{M^2 - a^2})^3}{M^3} \quad (5.2)$$

Until now, the deformation parameters had been allowed to take only positive values. These constraints show that this was potentially a naive restriction, thus the code should be adapted to allow for negative values with these constraints. This would ideally be done through `SpectralFitting.jl`'s parameter patching functionality or equivalent however if that is possible remains to be seen. For now, the lower bound for the deformation parameters has been set to -1, which is allowed no matter how large a gets.



6 Developing and Migrating to a Table Model

6.1 Introduction

Due to having to render the line profile through Gradus upon every iteration of the fit, the model took a significant amount of time to fit to the data. Thus other methods were looked into and it was decided to develop a table model, allowing for the pre-computation of the line profiles. The computed profiles could then be interpolated between for a given set of parameters, significantly speeding up computation. This would of course decrease the accuracy of the model, however the increased efficiency would likely be worth that loss.

The structure for XSPEC table models is well documented at [XSPEC Table Models](#) [4], thus the model was constructed using this specification. `SpectralFitting.jl` also has compatibility with this model format, which made integration more viable (see [here](#) for more [5]).

6.2 Data Generation

The spectra for the model were simply generated by looping through a defined parameter space, with a given number of values for each of the 5 parameters. The values selected were defined as follows

Parameter	Lower Bound	Upper Bound	Number
a	-0.998	0.998	7
h	1.5	30	10
θ	5	85	10
α_{13}	0	50	7
ϵ_3	0	30	10

This gave a total of 49,000 combinations to loop through. Seven values were used for both a and α_{13} as they had the smallest, and most predictable effect on the line profile, respectively. In a later iteration of this model it would be ideal to have a larger parameter space to reduce the inaccuracies of interpolation, however for a test case it was necessary to keep compute times down. The line profiles were stored in columns of a CSV with the header containing the parameter information.

Certain parameter combinations would fail to compute, and therefore a log file was kept, noting the combination and when it failed. This allowed the computations to run over night without risk of an error. The final compute time for this data generation was 16 hours. 3,500 of the 49,000 did not compute correctly, thus there is 45,500 datapoints in the grid.

6.3 Loading into FITS Format

This data was then loaded into a FITS file following the XSPEC model format using the python wrapper for the `heaspc` C library, `heasoftpy` (see the [documentation](#) for more information). The model specification requires that the grid be rectilinear [4], thus it was not possible to simply disregard the failed runs, and the spectra for these combinations was set to be zeroed out. This was determined to be acceptable as the majority of the

failed runs had large α_{13} or ϵ_3 , which from earlier testing was noted to be unlikely during fitting, thus the failed combinations would be unlikely to be used for interpolation.

6.4 The Model

To create a model in `SpectralFitting` using the table, the data was loaded into a `TableModelData` struct. This could then be wrapped using `TableModelInterpolation` such that it can work with the `TableModelInterpolation` function. This functionality made it trivial to set up the interpolation through the parameter grid by simply calling

```
SpectralFitting.interpolate_table!(table, a, h, , 13, 3)
```

This returned a spectrum with 1000 datapoints, which could then be interpolated over the domain of the input data. With this architecture it was simple to fit the data as previous.

6.4.1 Performance

Using the nu80402315002 dataset as before, the new model yielded the following results.




Figure 6.1: The table model result for the Kerr metric




Figure 6.2: The table model result for the Johannsen metric

Comparing to the previous fits, the χ^2 value for both of these fits were surprisingly improved. The Kerr fit was completed by freezing the deformation parameters at 0.

These fits took only 30 seconds to fit and plot both metrics, which was significantly faster than previous which could take an hour or more for a single fit.

6.5 Next Steps

After review of the log file, it was found that a lot of the cases where the deformation parameters were zero failed. This is significant as it implies that many of the Kerr model interpolations are inaccurate due to the “correction” method used (setting missing values to zero). Using code to parse the unique values for each of the 5 parameters, it was found that every instance of failure had $h = 1.5$. It was thus decided that upon the next iteration of the model, the lower bound of h would be set to 3. This value was determined after some brief testing with the failed runs, manually setting $h = 3$. $h = 2$ was also attempted, however this similarly caused errors.

Building on what was discussed in the previous chapter regarding the lower bounds for the parameters, it was also tested whether a value of -1 computed correctly for both of the deformation parameters. This was successful and thus a partial implementation of these bounds can be completed.

Further testing was also conducted in the parameter space of a where it was found that the line profile does not change with the sign of a . This will allow for much more resolution without increasing the length of computation.

To use the negative deformation parameters, their allowed values must be reduced as it is necessary to have confidence in the deformation = 0 (Kerr) cases. They will therefore take the values $\alpha_{13}, \epsilon_3 \in [-1, 10], \mathbb{Z}$. The other variables will take $h \in [3, 15]$, $a \in [0, 0.998]$ and $\theta \in [5, 85]$. Thus 7 values will be taken for a , 12 for the deformation parameters, and 9 for h , θ , giving a total size of 81648 for the table.

The final iteration that needs to be made upon the data is to increase the domain, as the model is currently forced to extrapolate causing unexpected issues, likely forcing the fitting to take worse values to avoid these errors.

All of these iterations were implemented and another dataset computed. At the time of writing these computations are still ongoing and therefore time was taken to clean up the plotting and do further reading.

Functionality was then added to plot a contour plot of the fit statistic (χ^2) for a pair of fit parameters. The fits are however not currently accurate enough to determine the functionality of this plot, thus this remains to be used.

6.6 Current Issues 2026/07/03

The table model currently does not contain the full set of spectra for negative values of the deformation parameters. For high spin parameters, the code will sometimes fail to compute the line profile citing either one of two errors:

```
ErrorException("No boundaries for minimization could be determined. It is likely this configuration
```

or

```
DomainError([VALUE], "sqrt was called with a negative real argument but will only return a complex r
```

This occurs when attempting to use negative values below ~ -0.4 for either, or both, of the deformation parameters. This is likely either a bug, or an artefact of the Johannsen metric. The only currently known constraints on allowed values of α_{13} and ϵ_3 are those stated in Equations Equation 5.1 and Equation 5.2, which all values used satisfy. The current working theory is thus that this may be a bug with `Gradus.jl` however this requires further research.

There also exists the restriction that $|a| < 1$ to avoid a naked singularity [6], however this is once again satisfied by the used values.

The table model is currently capable of extrapolating, which provides reasonable estimates of the line profile for negative values (see Figure 6.3) however it would not be preferable to rely on extrapolation to constrain the parameters.




Figure 6.3

7 Notes for the Reduction of XRISM Data

With HEASOFT and CalDB installed and instantiated:

- Data will have observation ID as top directory, e.g. 000125000/
- Pipeline scripts must be run in a directory that contains this directory

7.1 Pipeline Scripts

For all following commands replace 000125000 with the obsid for your data. All commands here are run starting from the 000125000/ directory.

7.1.1 Reprocessing:

In the directory above data directory run the following for the resolve data

```
# Run the pipeline
xapipeline indir=000125000 outdir=000125000_rsl_reproc steminputs=xa000125000 stemoutputs=DEFAULT er

# Pack and cleanup logs
7z a 000125000_rsl_logs.7z *.log
rm *.log

# Enter the directory, remove duplicate files and gzip
cd 000125000_rsl_reproc
rm *.hk
gzip *

cd ..
```

Then for the Xtend data

```
# Run the pipeline
xapipeline indir=000125000 outdir=000125000_xtd_reproc steminputs=xa000125000 stemoutputs=DEFAULT er

# Pack and cleanup logs
7z a 000125000_xtd_logs.7z *.log
rm *.log

# Enter the directory, remove duplicate files and gzip
cd 000125000_xtd_reproc
rm *.hk
gzip *

cd ..
```


7.1.2 Analysis Setup

Then make a directory for the analysis and copy the following files

```
mkdir analysis

cp 000125000_xtd_reproc/*_cl.evt.gz analysis/
cp 000125000_rsl_reproc/*_cl.evt.gz analysis/
cp 000125000_rsl_reproc/xa000125000.ehk.gz 000125000_rsl_reproc/xa000125000rsl_px1000_exp.gti.gz 000125000_rsl_reproc/xa000125000rsl_px1000_exp.gti.gz
```

7.1.3 PSP Limit Checks

Run PSP limit check to see if the unscreened count rate exceeds the psp limit.

```
rs1bratios infile=xa000125000rsl_p0px1000_uf.evt.gz filetype=uf outroot=rs1000125000_ufbr
```

This gives a info on if quadrants exceed 50 cts s^{-1} , which makes the data unusable. Next run rs1bratios on the cleaned resolve data:

```
rs1bratios infile=xa000125000rsl_p0px1000_cl.evt.gz filetype=cl outroot=rs1000125000_clbr lcbn=128
```

This gives the same info but for the cleaned data and outputs the following which will be used later

- rs1000125000_clbr_2keVto12keV_hist.fits: Observed vs. predicted branching ratios in the energy range specified by the “eband” parameter
- rs1000125000_clbr_full_lc.fits: Full array light curves by grade
- rs1000125000_clbr_pix_lc.fits: Per-pixel light curves by grade
- rs1000125000_clbr_gradeFrac_lc.fits: Grade fractions versus time
- rs1000125000_clbr_2keVto12keV_pixMap.img: 20 images of counts by grade, and grade fractions for two energy bands

7.2 Removing Anomalous Pixels from Xtend Data

All following commands are run from the 000125000_xtd_reproc/ directory

First generate region files using ds9, viewing the image in DET coordinates with the command.

```
ds9 "xa000125000xtd_p031100010_cl.evt.gz[bin=detx,dety]"
```




Figure 7.1: The regions used.

Three regions were selected and labelled as follows:


- White: 1_bottom.reg
- Green: 1_src.reg
- Cyan: 1_top.reg

```


punlearn xtdpixclip
xtdpixclip evtinfile="xa000125000xtd_p031100010_cl.evt.gz" outroot="xa000125000xtd_p031100010_cl" pr

```


The above is used to determine the thresholds to use as cutoffs for anomolous pixels. The output histograms are shown below



(a) The bottom region histogram



(a) The source region histogram




(a) The top region histogram

from these histograms it is possible to select the thresholds to cut off anomolous data. For the two non-source regions a threshold of 8 was selected, and for the source region 1000 was selected. With the regions selected run


```

punlearn xtdpixclip
xtdpixclip evtinfile="xa000125000xtd_p031100010_cl.evt.gz" outroot="xa000125000xtd_p031100010_cl" pr


```



(a) The Raw image



(a) The cleaned image



(a) The filtered out pixels

7.3 Proximity and Rise time screening


- Proximity screening: Removing frame events based on the unlikelihood of events happening too close together in a given time interval. This must be balanced against the loss of counts.
- Rise-time screening:

The effect of both can be previewed through the FITS files

```
rs1000125000_clbr_prox_Hp_spectra.fits  
rs1000125000_clbr_rise_Hp_spectra.fits
```

These were examined using the following julia script

```
using FITSFiles  
using Plots  
  
file = fits("/home/brad/Downloads/000125000_rsl_reproc/rs1000125000_clbr_rise_Hp_spectra.fits")  
display(plot(file[3].data.PI, file[3].data.DIFF; xlabel="PI Channel", ylabel="Difference (counts)",  
file = fits("/home/brad/Downloads/000125000_rsl_reproc/rs1000125000_clbr_prox_Hp_spectra.fits")  
display(plot(file[3].data.PI, file[3].data.DIFF; xlabel="PI Channel", ylabel="Difference (counts)",
```



The losses for each were determined to be worth it for this dataset, thus the both filters were applied with the command:

```
ftcopy infile="xa000125000rsl_p0px1000_cl.evt[EVENTS] [(PI>=600) && (((RISE_TIME+0.00075*DERIV_MAX,
```

To skip proximity filtering you can omit `&&STATUS[4]==b0`.

7.4 Extract Broadband Images and Establish Source Center Coordinates

- Check the value of the EXPOSURE keyword in the cleaned event files, if the exposure time is unexpectedly low then consult section 9 of the [XRISM QSG](#).
- You should now have the following files
 - xa000125000rsl_p0px1000_cl2.evt (Resolve) and
 - xa000125000xtd_p031100010_cl.evt or xa000125000xtd_p031100010_xpc_clnevt.fits (Xtend)

Start XSELECT and make a 2-10 keV resolve image in DET coordinates as:

```
terminal> xselect
xselect> read eve xa000125000rsl_p0px1000_cl2.evt .
xselect> set image DET
xselect> filter pha_cutoff 4000 20000
xselect> extr image
xselect> save image xa000125000rsl_p0px1000_detimg.fits
```

The centre should be close to (3.5, 3.5), we only need to estimate this approximately as it is used as a sanity check against the RA and DEC coordinates. For most standard observations we can simply assume this centre coordinate.

Next copy over a resolve region file

```
cp $HEADAS/refdata/region_RSL_det.reg .
```


and make another region file `exclude_calsource.reg` containing

```
physical
-circle(920.0,1530.0,92.0)
-circle(919.0,271.0,91.0)
```

Next make a 0.5-10 keV Xtend image in DET coordinates

```
xselect> clear all
xselect> read eve xa000125000xtd_p031100010_xpc_clnevt.fits.gz .
xselect> set image DET
xselect> filter region exclude_calsources.reg
xselect> filter pha_cutoff 83 1667
xselect> extr image
xselect> save image xa000125000xtd_p031100010_detimg.fits
```

Now we can estimate the centroid of the source in DET coordinates using `ds9`, it is recommended to use a rectangle of length 5 arcmin and width as wide as possible for 1/8 Window mode images. For full window mode it is recommended to use a circle of radius 2.5 arcmin. The regions selected here are shown below



and were saved as

```
region_000125000xtd_src.reg  
region_000125000xtd_bgd.reg
```

References

- [1] A. C. Fabian *et al.*, Publications of the Astronomical Society of the Pacific **112**, 1145 (2000).
- [2] F. A. Harrison *et al.*, American Physical Journal **770**, 103 (2013).
- [3] T. Johannsen, Physical Review D **88**, (2013).
- [4] K. A. Arnaud, *The File Format for XSPEC Table Models* (NASA, Greenbelt, MD20771, 2021).
- [5] F. Baker and A. Young, Zenodo (2025).
- [6] A. Tripathi *et al.*, Phys. Rev. D **99**, 083001 (2019).